

Ablation study: capacity, data size, and training budget

Hypothesis: larger models should improve likelihood, but multi-ticker data and longer training must be judged with both token metrics and rollout facts.

Run	Backtest loss	Price-depth L1
one/micro	3.060 	0.632 
one/small	2.114 	0.146 
two/micro	4.771 	1.229 
two/small	4.943 	0.141 
small_100	4.166 	0.801 

Conclusion: small beats micro on the one-ticker slice; adding a second ticker is harder; the 100-step follow-up improves loss but worsens price-depth L1, so selection must monitor both open-loop and rollout metrics.